

Case assignment

Fjord Global Investments has grown from 5 subsidiaries to around 100 in under a decade, and our governance data is scattered across spreadsheets, emails, and letters from local agents in 18 jurisdictions. Your job is to bring order to this chaos. Ingest our messy data, surface the risks we can't currently see, and show us how AI and smart tooling can give a small legal team superpowers. Build something you'd be proud to demo to our General Counsel on Thursday.

Requirements

1. **Frontend** — Build a UI in React or Angular. The core requirement is that users can trigger a digest fetch and view the results, but we encourage you to think beyond the minimum — filtering, sorting, status indicators, or anything else that makes the tool genuinely useful. Show us what you'd build if this were a real product.
2. **Backend** — Build a FastAPI service that calls an LLM API (any provider or open-source model) to fetch, summarise, and filter content, and to produce action recommendations for each item surfaced.
3. **How you store and structure your data is entirely up to you** — use whatever approach fits your design, whether that's a local file, SQLite, PostgreSQL, Firebase, or something else. We're more interested in your reasoning than your infrastructure choices.

Using AI

We encourage you to use AI tools throughout this assignment — for coding, for structuring your thinking, or for anything else where it adds value. We will reimburse up to USD 30 in LLM API costs; just keep receipts, and you can request reimbursement after the interview is completed. That said, we are as interested in your process as your output. Be prepared to walk us through how you used AI tools, where they helped, and where you had to step in yourself. We will spend part of the interview exploring this together.

Submission

Please upload your code to a public GitHub repository and send the link to tobias.hyldmo@nbim.no, martin.espeland@nbim.no and guh@nbim.no at least 24 hours before your interview.

Presentation: Show the application in the meeting for 5–10 minutes, followed by a short discussion of your process and design choices. Remember to bring your laptop so you can run the application locally on your machine.